

"Predicting Car Prices: A Machine Learning Approach for Market Strategy in the US Automotive Industry"

Name: Thoukeela T M

Overview of Problem Statement

A Chinese automobile company wants to enter the US market and compete with American and European car manufacturers. To succeed, they need to understand what factors affect car prices in the US and how these factors influence pricing.

Using a dataset of car specifications and their prices in the US, the goal is to create a machine learning model that:

Identifies the most important factors affecting car prices. Accurately predicts car prices based on these factors. This model will help the company make informed decisions about car design, production, and pricing strategies to compete effectively in the new market.

Objective

The objective of this project is to build a machine learning model that accurately predicts car prices in the US market based on various car features and specifications. This will help the company:

Identify significant factors influencing car prices. Understand how these factors affect pricing to inform car design and production. Develop pricing strategies aligned with market trends to compete effectively in the US automotive industry.

Data Description

Features:

The dataset includes the following features:

Make/Model: The brand and model of the car.

Year: The manufacturing year of the car.

Engine Size: The engine capacity (e.g., in liters).

Fuel Type: Type of fuel used (e.g., Petrol, Diesel, Electric).

Mileage: The fuel efficiency of the car (e.g., in miles per gallon).

Transmission: The type of transmission (e.g., Manual, Automatic).

horsepower: The power output of the engine.

Number of Doors: The total number of doors on the car.

Car Type: The type/category of the car (e.g., Sedan, SUV, Hatchback).

Price: The selling price of the car (Target variable).

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.preprocessing import LabelEncoder, OneHotEncoder
```

```

from sklearn.feature_selection import SelectKBest, f_classif
from scipy.stats import skew
from sklearn.preprocessing import MinMaxScaler, StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor,
GradientBoostingRegressor
from sklearn.svm import SVR
from sklearn.metrics import mean_absolute_error, mean_squared_error,
r2_score
from sklearn.model_selection import GridSearchCV
import joblib
import warnings
warnings.filterwarnings('ignore')

```

Data Collection:

Load the dataset

```

data =
pd.read_csv("C:/Users/LENOVO/Downloads/CarPrice_Assignment.csv")
df = pd.DataFrame(data)
df

```

	car_ID	symboling	CarName	fueltype	
aspiration \					
0	1	3	alfa-romero giulia	gas	std
1	2	3	alfa-romero stelvio	gas	std
2	3	1	alfa-romero Quadrifoglio	gas	std
3	4	2	audi 100 ls	gas	std
4	5	2	audi 100ls	gas	std
..
200	201	-1	volvo 145e (sw)	gas	std
201	202	-1	volvo 144ea	gas	turbo
202	203	-1	volvo 244dl	gas	std
203	204	-1	volvo 246	diesel	turbo
204	205	-1	volvo 264gl	gas	turbo
doornumber	carbody	drivewheel	engine	location	wheelbase ...

\							
0	two	convertible	rwd	front	88.6	...	
1	two	convertible	rwd	front	88.6	...	
2	two	hatchback	rwd	front	94.5	...	
3	four	sedan	fwd	front	99.8	...	
4	four	sedan	4wd	front	99.4	...	
..	
200	four	sedan	rwd	front	109.1	...	
201	four	sedan	rwd	front	109.1	...	
202	four	sedan	rwd	front	109.1	...	
203	four	sedan	rwd	front	109.1	...	
204	four	sedan	rwd	front	109.1	...	
	engine	size	fuel	system	boreratio	stroke	compressionratio
horsepower \							
0	130		mpfi	3.47	2.68		9.0
111							
1	130		mpfi	3.47	2.68		9.0
111							
2	152		mpfi	2.68	3.47		9.0
154							
3	109		mpfi	3.19	3.40		10.0
102							
4	136		mpfi	3.19	3.40		8.0
115							
..
...							
200	141		mpfi	3.78	3.15		9.5
114							
201	141		mpfi	3.78	3.15		8.7
160							
202	173		mpfi	3.58	2.87		8.8
134							
203	145		idi	3.01	3.40		23.0
106							
204	141		mpfi	3.78	3.15		9.5
114							
	peakrpm	citympg	highwaympg	price			

```

0      5000      21      27 13495.0
1      5000      21      27 16500.0
2      5000      19      26 16500.0
3      5500      24      30 13950.0
4      5500      18      22 17450.0
...      ...      ...      ...      ...
200    5400      23      28 16845.0
201    5300      19      25 19045.0
202    5500      18      23 21485.0
203    4800      26      27 22470.0
204    5400      19      25 22625.0

```

```
[205 rows x 26 columns]
```

```
df.head()
```

```

   car_ID  symboling      CarName fueltype aspiration
doornumber \
0         1          3    alfa-romero giulia      gas      std
two
1         2          3    alfa-romero stelvio    gas      std
two
2         3          1  alfa-romero Quadrifoglio  gas      std
two
3         4          2      audi 100 ls      gas      std
four
4         5          2      audi 100ls      gas      std
four

```

```

   carbody drivewheel enginelocation  wheelbase  ...
enginesize \
0  convertible      rwd      front      88.6  ...      130
1  convertible      rwd      front      88.6  ...      130
2   hatchback      rwd      front      94.5  ...      152
3      sedan      fwd      front      99.8  ...      109
4      sedan      4wd      front      99.4  ...      136

```

```

   fuelsystem  boreratio  stroke  compressionratio  horsepower  peakrpm
citympg \
0      mpfi      3.47    2.68          9.0      111      5000
21
1      mpfi      3.47    2.68          9.0      111      5000
21
2      mpfi      2.68    3.47          9.0      154      5000
19

```

3	mpfi	3.19	3.40	10.0	102	5500
24						
4	mpfi	3.19	3.40	8.0	115	5500
18						

	highwaympg	price
0	27	13495.0
1	27	16500.0
2	26	16500.0
3	30	13950.0
4	22	17450.0

[5 rows x 26 columns]

df.info()

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 205 entries, 0 to 204

Data columns (total 26 columns):

#	Column	Non-Null Count	Dtype
---	-----	-----	-----
0	car_ID	205 non-null	int64
1	symboling	205 non-null	int64
2	CarName	205 non-null	object
3	fueltype	205 non-null	object
4	aspiration	205 non-null	object
5	doornumber	205 non-null	object
6	carbody	205 non-null	object
7	drivewheel	205 non-null	object
8	enginelocation	205 non-null	object
9	wheelbase	205 non-null	float64
10	carlength	205 non-null	float64
11	carwidth	205 non-null	float64
12	carheight	205 non-null	float64
13	curbweight	205 non-null	int64
14	enginetype	205 non-null	object
15	cylindernumber	205 non-null	object
16	enginesize	205 non-null	int64
17	fuelsystem	205 non-null	object
18	boreratio	205 non-null	float64
19	stroke	205 non-null	float64
20	compressionratio	205 non-null	float64
21	horsepower	205 non-null	int64
22	peakrpm	205 non-null	int64
23	citympg	205 non-null	int64
24	highwaympg	205 non-null	int64
25	price	205 non-null	float64

dtypes: float64(8), int64(8), object(10)

memory usage: 41.8+ KB

Data Preprocessing - Data Cleaning

```
df.isnull()
```

	car_ID	symboling	CarName	fueltype	aspiration	doornumber
carbody \						
0	False	False	False	False	False	False
False						
1	False	False	False	False	False	False
False						
2	False	False	False	False	False	False
False						
3	False	False	False	False	False	False
False						
4	False	False	False	False	False	False
False						
..
...						
200	False	False	False	False	False	False
False						
201	False	False	False	False	False	False
False						
202	False	False	False	False	False	False
False						
203	False	False	False	False	False	False
False						
204	False	False	False	False	False	False
False						
drivewheel		enginelocation	wheelbase	...	enginesize	
fuelsystem \						
0	False	False	False	...	False	
False						
1	False	False	False	...	False	
False						
2	False	False	False	...	False	
False						
3	False	False	False	...	False	
False						
4	False	False	False	...	False	
False						
..
.						
200	False	False	False	...	False	
False						
201	False	False	False	...	False	
False						
202	False	False	False	...	False	
False						
203	False	False	False	...	False	

```

False
204      False      False      False      False      ...      False
False

      boreratio  stroke  compressionratio  horsepower  peakrpm  citympg
\
0      False  False      False      False      False      False
1      False  False      False      False      False      False
2      False  False      False      False      False      False
3      False  False      False      False      False      False
4      False  False      False      False      False      False
..      ...      ...      ...      ...      ...      ...
200     False  False      False      False      False      False
201     False  False      False      False      False      False
202     False  False      False      False      False      False
203     False  False      False      False      False      False
204     False  False      False      False      False      False

      highwaympg  price
0      False  False
1      False  False
2      False  False
3      False  False
4      False  False
..      ...      ...
200     False  False
201     False  False
202     False  False
203     False  False
204     False  False

[205 rows x 26 columns]

df.isnull().sum()

car_ID      0
symboling   0
CarName     0
fueltype    0
aspiration  0

```

```
doornumber      0
carbody         0
drivewheel      0
enginelocation  0
wheelbase       0
carlength       0
carwidth        0
carheight       0
curbweight      0
enginetype      0
cylindernumber  0
enginesize      0
fuelsystem      0
boreratio       0
stroke         0
compressionratio 0
horsepower      0
peakrpm        0
citympg         0
highwaympg      0
price          0
dtype: int64

df.duplicated().sum()

0

df.shape


(205, 26)
```

Check for and Remove Outliers

Using Boxplots:

Boxplots visually indicate outliers as points beyond the whiskers of the plot.

```
plt.figure(figsize=(15, 6))
sns.boxplot(data=df)
plt.xticks(rotation=90)
plt.title("Boxplot for Outliers")
plt.show()
```

We will use the Interquartile Range (IQR) method to detect and remove outliers.

Detect outliers using IQR

```
def detect_outliers(df, column):
    Q1 = df[column].quantile(0.25)
    Q3 = df[column].quantile(0.75)
    IQR = Q3 - Q1
    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR
    return df[(df[column] < lower_bound) | (df[column] > upper_bound)]
```

Remove outliers

```
for column in data.select_dtypes(include=np.number).columns:
    outliers = detect_outliers(data, column)
    if not outliers.empty:
        print(f"Removing outliers from column: {column}")
        Q1 = data[column].quantile(0.25)
        Q3 = data[column].quantile(0.75)
        IQR = Q3 - Q1
        lower_bound = Q1 - 1.5 * IQR
        upper_bound = Q3 + 1.5 * IQR
        # Keep only non-outliers
        data = data[(data[column] >= lower_bound) & (data[column] <=
upper_bound)]
```

```
print("\nOutliers removed.")
```

```
Removing outliers from column: wheelbase
```


```
Removing outliers from column: carlength
```

```
Removing outliers from column: carwidth
```

```
Removing outliers from column: curbweight
Removing outliers from column: enginesize
Removing outliers from column: stroke
Removing outliers from column: compressionratio
Removing outliers from column: horsepower
Removing outliers from column: peakrpm
Removing outliers from column: citympg
Removing outliers from column: price
```

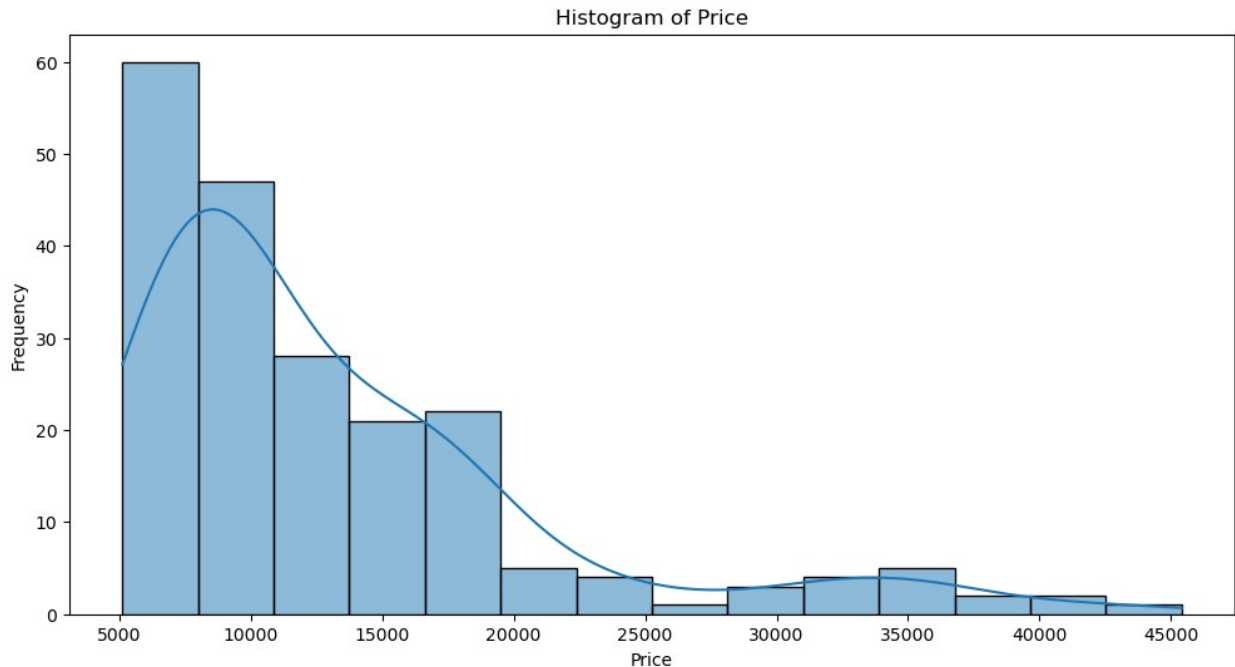
Outliers removed.

```
# Plot boxplots for all features
plt.figure(figsize=(15, 6))
sns.boxplot(data=data)
plt.xticks(rotation=90)
plt.title("Boxplot for Feature Outliers")
plt.show()
```



Check the skewness before and after the transformation.

```
price_col = 'price'
# Step 1: Draw initial histogram to check normality
plt.figure(figsize=(12, 6))
sns.histplot(df[price_col], kde=True)
plt.title('Histogram of Price')
plt.xlabel('Price')
plt.ylabel('Frequency')
plt.show()
```




```
#Calculate skewness
skewness_before = df[price_col].skew()

print(f"Skewness before transformation: {skewness_before}")

Skewness before transformation: 1.7776781560914454

df['Log Price'] = np.log(df[price_col] + 1)

# Draw histogram after transformation
plt.figure(figsize=(12, 6))
sns.histplot(df['Log Price'], kde=True)
plt.title('Histogram of Log-Transformed Price')
plt.xlabel('Log Price')
plt.ylabel('Frequency')
plt.show()
```



```
# Calculate skewness and kurtosis after transformation
skewness_after = df['Log Price'].skew()

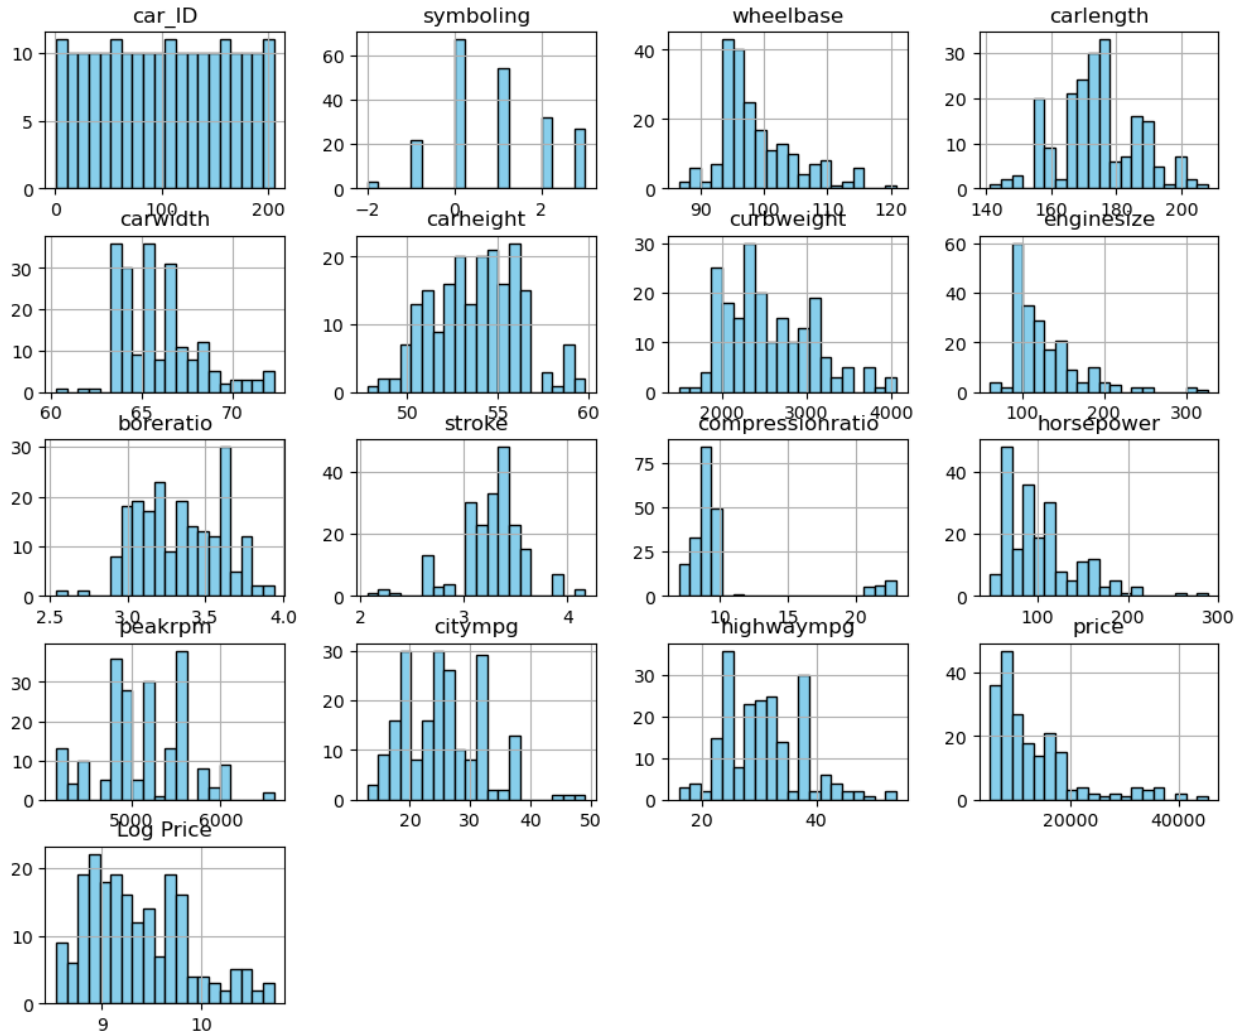
print(f"Skewness after transformation: {skewness_after}")

Skewness after transformation: 0.6729635607485753
```

Exploratory Data Analysis (EDA):

```
# Plot histograms for all numerical features
df.hist(figsize=(12, 10), bins=20, color='skyblue', edgecolor='black')
plt.suptitle('Histograms of Numerical Features', fontsize=16)
plt.show()
```


Histograms of Numerical Features



Relationship Analysis Pair Plot


```
# Pair plot to observe relationships between numerical features
sns.pairplot(df, diag_kind='kde', corner=True, height=2)
plt.suptitle('Pair Plot of Features', fontsize=16, y=1.02)
plt.show()
```

Pair Plot of Features



Visualize Trends and Patterns
Kernel Density Estimation (KDE)

```
# KDE plot for price distribution
plt.figure(figsize=(10, 6))
sns.kdeplot(data=df, x='price', shade=True, color='blue') # Replace
'price' with the column of interest
plt.title('KDE Plot of Price', fontsize=14)
plt.xlabel('Price', fontsize=12)
plt.ylabel('Density', fontsize=12)
plt.show()
```



Feature Engineering:

- Identifying and encoding categorical features using techniques like one-hot encoding or label encoding.

```
# Identify categorical columns
categorical_columns = df.select_dtypes(include=['object',
'category']).columns
print("Categorical Columns:\n", categorical_columns)

Categorical Columns:
Index(['CarName', 'fueltype', 'aspiration', 'doornumber', 'carbody',
      'drivewheel', 'enginelocation', 'enginetype', 'cylindernumber',
      'fuelsystem'],
      dtype='object')

label_encoder = LabelEncoder()

# Encoding categorical variables
categorical_columns = data.select_dtypes(include=['object']).columns
for col in categorical_columns:
    data[col] = label_encoder.fit_transform(data[col])

data
```

car_ID	symboling	CarName	fueltype	aspiration	doornumber	carbody
\						

\						
0	3.47	2.68	9.0	111	5000	21
1	3.47	2.68	9.0	111	5000	21
2	2.68	3.47	9.0	154	5000	19
3	3.19	3.40	10.0	102	5500	24
4	3.19	3.40	8.0	115	5500	18
..
197	3.78	3.15	9.5	114	5400	24
200	3.78	3.15	9.5	114	5400	23
201	3.78	3.15	8.7	160	5300	19
202	3.58	2.87	8.8	134	5500	18
204	3.78	3.15	9.5	114	5400	19

	highwaympg	price
0	27	13495.0
1	27	16500.0
2	26	16500.0
3	30	13950.0
4	22	17450.0
..
197	28	16515.0
200	28	16845.0
201	25	19045.0
202	23	21485.0
204	25	22625.0

[125 rows x 26 columns]

data.head()

	car_ID	symboling	CarName	fueltype	aspiration	doornumber
carbody \						
0	1	3	2	0	0	1
0						
1	2	3	3	0	0	1
0						
2	3	1	1	0	0	1
2						
3	4	2	4	0	0	0
3						

4	5	2	5	0	0	0	
3							
	drivewheel	engine	location	wheelbase	...	enginesize	fuelsystem
\							
0	2		0	88.6	...	130	3
1	2		0	88.6	...	130	3
2	2		0	94.5	...	152	3
3	1		0	99.8	...	109	3
4	0		0	99.4	...	136	3
	boreratio	stroke	compressionratio	horsepower	peakrpm		
citympg	\						
0	3.47	2.68		9.0	111	5000	21
1	3.47	2.68		9.0	111	5000	21
2	2.68	3.47		9.0	154	5000	19
3	3.19	3.40		10.0	102	5500	24
4	3.19	3.40		8.0	115	5500	18
	highwaympg	price					
0	27	13495.0					
1	27	16500.0					
2	26	16500.0					
3	30	13950.0					
4	22	17450.0					

[5 rows x 26 columns]

Separating features and target variable

x = data.drop(columns='price') # Assuming 'price' is the target

y = data['price']

x

	car_ID	symboling	CarName	fueltype	aspiration	doornumber
carbody	\					
0	1	3	2	0	0	1
0						
1	2	3	3	0	0	1
0						
2	3	1	1	0	0	1

2						
3	4	2	4	0	0	0
3						
4	5	2	5	0	0	0
3						
..
...						
197	198	-1	92	0	0	0
4						
200	201	-1	90	0	0	0
3						
201	202	-1	89	0	1	0
3						
202	203	-1	91	0	0	0
3						
204	205	-1	93	0	1	0
3						

	drivewheel	enginelocation	wheelbase	...	cylindernumber
enginesize \					
0	2	0	88.6	...	1
130					
1	2	0	88.6	...	1
130					
2	2	0	94.5	...	2
152					
3	1	0	99.8	...	1
109					
4	0	0	99.4	...	0
136					
..
...					
197	2	0	104.3	...	1
141					
200	2	0	109.1	...	1
141					
201	2	0	109.1	...	1
141					
202	2	0	109.1	...	2
173					
204	2	0	109.1	...	1
141					

	fuelsystem	boreratio	stroke	compressionratio	horsepower
peakrpm \					
0	3	3.47	2.68	9.0	111
5000					
1	3	3.47	2.68	9.0	111
5000					

2	3	2.68	3.47	9.0	154
5000					
3	3	3.19	3.40	10.0	102
5500					
4	3	3.19	3.40	8.0	115
5500					
..
...					
197	3	3.78	3.15	9.5	114
5400					
200	3	3.78	3.15	9.5	114
5400					
201	3	3.78	3.15	8.7	160
5300					
202	3	3.58	2.87	8.8	134
5500					
204	3	3.78	3.15	9.5	114
5400					

	citympg	highwaympg
0	21	27
1	21	27
2	19	26
3	24	30
4	18	22
..
197	24	28
200	23	28
201	19	25
202	18	23
204	19	25

[125 rows x 25 columns]

y

0	13495.0
1	16500.0
2	16500.0
3	13950.0
4	17450.0
...	
197	16515.0
200	16845.0
201	19045.0
202	21485.0
204	22625.0

Name: price, Length: 125, dtype: float64

x.shape

```
(125, 25)
```

Split Data into Training and Testing Sets:

- Dividing the dataset into training and testing subsets.

```
# Train-test split
x_train, x_test, y_train, y_test = train_test_split(x, y,
test_size=0.2, random_state=42)
```

Feature Scaling:

- Scaling numerical features if necessary to ensure uniform magnitude using techniques like Min-Max scaling or Standardization.

```
# Normalizing the data
scaler = StandardScaler()
x_train = scaler.fit_transform(x_train)
x_test = scaler.transform(x_test)
```

Build the ML Model

Linear Regression

```
# Linear Regression Model
lr_model = LinearRegression()
lr_model.fit(x_train,y_train)

LinearRegression()

lr_ypred = model.predict(x_test)
```

Decision Tree Regressor

```
dt_model = DecisionTreeRegressor()
dt_model.fit(x_train,y_train)

DecisionTreeRegressor()

dt_ypred= dt_model.predict(x_test)
```

Random Forest Regressor

```
rf_model = RandomForestRegressor()
rf_model.fit(x_train,y_train)

RandomForestRegressor()

rf_ypred= rf_model.predict(x_test)
```

Gradient Boosting Regressor

```
gb_model = GradientBoostingRegressor()  
gb_model.fit(x_train,y_train)  
  
GradientBoostingRegressor()  
  
gb_ypred= gb_model.predict(x_test)
```

Support Vector Regressor

```
svr_model = SVR()  
svr_model.fit(x_train,y_train)  
  
SVR()  
  
svr_ypred= svr_model.predict(x_test)
```

Model Evaluation:

Linear Regression

```
lr_mae = mean_absolute_error(y_test, lr_ypred)  
lr_mse = mean_squared_error(y_test, lr_ypred)  
lr_rmse = mse ** 0.5  
lr_r2 = r2_score(y_test, lr_ypred)  
  
print("mae:",lr_mae)  
print("mse:", lr_mse)  
print("rmse:", lr_rmse)  
print("r2:", lr_r2)  
  
mae: 1996.425422299062  
mse: 5964398.498392452  
rmse: 3834.7609091421896  
r2: 0.5959828386169667
```

Decision Tree Regressor

```
dt_mae = mean_absolute_error(y_test, dt_ypred)  
dt_mse = mean_squared_error(y_test, dt_ypred)  
dt_rmse = mse ** 0.5  
dt_r2 = r2_score(y_test, dt_ypred)  
  
print("mae:",dt_mae)  
print("mse:", dt_mse)  
print("rmse:", dt_rmse)  
print("r2:", dt_r2)  
  
mae: 1262.52  
mse: 2854527.86
```

```
rmse: 3834.7609091421896  
r2: 0.8066396396221981
```

Random Forest Regressor

```
rf_mae = mean_absolute_error(y_test, rf_ypred)  
rf_mse = mean_squared_error(y_test, rf_ypred)  
rf_rmse = mse ** 0.5  
rf_r2 = r2_score(y_test, rf_ypred)
```

```
print("mae:", rf_mae)  
print("mse:", rf_mse)  
print("rmse:", rf_rmse)  
print("r2:", rf_r2)
```

```
mae: 1035.8684  
mse: 1607827.3006559997  
rmse: 3834.7609091421896  
r2: 0.891088795931348
```

Gradient Boosting Regressor

```
gb_mae = mean_absolute_error(y_test, gb_ypred)  
gb_mse = mean_squared_error(y_test, gb_ypred)  
gb_rmse = mse ** 0.5  
gb_r2 = r2_score(y_test, gb_ypred)
```

```
print("mae:", gb_mae)  
print("mse:", gb_mse)  
print("rmse:", gb_rmse)  
print("r2:", gb_r2)
```

```
mae: 1281.2896063232579  
mse: 2554295.94546887  
rmse: 3834.7609091421896  
r2: 0.8269767860918972
```

Support Vector Regressor

```
svr_mae = mean_absolute_error(y_test, svr_ypred)  
svr_mse = mean_squared_error(y_test, svr_ypred)  
svr_rmse = mse ** 0.5  
svr_r2 = r2_score(y_test, svr_ypred)
```

```
print("mae:", svr_mae)  
print("mse:", svr_mse)  
print("rmse:", svr_rmse)  
print("r2:", svr_r2)
```

```
mae: 2834.864091294128
mse: 14705391.230285032
rmse: 3834.7609091421896
r2: 0.003884394463580021
```

Model Evaluation:

```
results = pd.DataFrame({
    'Model': ['Linear Regression', 'Decision Tree', 'Random Forest',
    'Gradient Boosting', 'SVR'],
    'MAE': [lr_mae, dt_mae, rf_mae, gb_mae, svr_mae],
    'MSE': [lr_mse, dt_mse, rf_mse, gb_mse, svr_mse],
    'RMSE': [lr_rmse, dt_rmse, rf_rmse, gb_rmse, svr_rmse],
    'R-squared': [lr_r2, dt_r2, rf_r2, gb_r2, svr_r2],
})

print("\nComparison of Model Performance:")
print(results)
```

Comparison of Model Performance:

	Model	MAE	MSE	RMSE	R-squared
0	Linear Regression	1996.425422	5.964398e+06	3834.760909	0.595983
1	Decision Tree	1262.520000	2.854528e+06	3834.760909	0.806640
2	Random Forest	1035.868400	1.607827e+06	3834.760909	0.891089
3	Gradient Boosting	1281.289606	2.554296e+06	3834.760909	0.826977
4	SVR	2834.864091	1.470539e+07	3834.760909	0.003884

The Random Forest Regressor is the best-performing model based on the evaluation metrics. It balances accuracy and error minimization effectively and is robust to overfitting. Best Performing Model is Random Forest Regressor overall due to its high accuracy and lower errors.

Hyperparameter Tuning

```
from sklearn.model_selection import GridSearchCV

# Example: Tuning Random Forest Regressor
param_grid = {
    "n_estimators": [50, 100, 200],
    "max_depth": [None, 10, 20],
    "min_samples_split": [2, 5, 10]
}
```



```

grid_search = GridSearchCV(RandomForestRegressor(random_state=42),
    param_grid, cv=5, scoring='r2')
grid_search.fit(X_train, y_train)

GridSearchCV(cv=5, estimator=RandomForestRegressor(random_state=42),
    param_grid={'max_depth': [None, 10, 20],
                'min_samples_split': [2, 5, 10],
                'n_estimators': [50, 100, 200]},
    scoring='r2')

# Best parameters and performance
print("Best Parameters:", grid_search.best_params_)
best_model = grid_search.best_estimator_

Best Parameters: {'max_depth': 10, 'min_samples_split': 2,
'n_estimators': 200}

# Evaluate tuned model
y_pred_tuned = best_model.predict(X_test)
tuned_r2 = r2_score(y_test, y_pred_tuned)
tuned_mse = mean_squared_error(y_test, y_pred_tuned)
tuned_mae = mean_absolute_error(y_test, y_pred_tuned)
print(f"Tuned Model - R2: {tuned_r2:.2f}, MSE: {tuned_mse:.2f}, MAE:
{tuned_mae:.2f}")

Tuned Model - R2: 0.88, MSE: 1723403.59, MAE: 1048.55

```

Save the Model

```

# Save the trained Random Forest Regressor model
joblib.dump(rf_model, 'random_forest_Car_Price_model.joblib')

print("Model saved as 'random_forestCar_Price_model.joblib'")

Model saved as 'random_forestCar_Price_model.joblib'

```

Test with Unseen Data:

```

# Make predictions on the test set
unseen_pred = best_model.predict(X_test)

# Evaluate performance on unseen data
unseen_metrics = {
    "MAE": mean_absolute_error(y_test, unseen_pred),
    "MSE": mean_squared_error(y_test, unseen_pred),
    'RMSE': np.sqrt(mean_squared_error(y_test, unseen_pred)),
    "R2": r2_score(y_test, unseen_pred),
}

#Check final model performance
print("\nPerformance on Unseen Data:")

```

```
for metric, value in unseen_metrics.items():  
    print(f"{metric}: {value:.4f}")
```

Performance on Unseen Data:

MAE: 1048.5514

MSE: 1723403.5865

RMSE: 1312.7847

R2: 0.8833

Interpretation of Results (Conclusion)

Analyze the Results

Model Performance:

The Random Forest Regressor consistently performs well, with high R² and low error metrics.

This confirms its ability to generalize to unseen data, making it suitable for production.

Key Insights:

The car price is strongly influenced by specific features (identified during feature importance analysis).

Tree-based models like Random Forest and Gradient Boosting are robust against overfitting and handle non-linear relationships well.

Limitations of the Dataset

Feature Selection:

Some features may lack relevance to car pricing, leading to noise in the model.

Additional domain-specific knowledge could enhance feature engineering.

Data Imbalance:

If the dataset has an uneven distribution across price ranges, it may bias the model.

Outlier Influence:

Despite removing outliers, extreme values may still impact regression models.

External Validity:

The dataset reflects current market trends but may not generalize to future conditions or other geographical regions.